




AI Pathways Technical Track

Capstone Reflection Journal



As you progress through Course 3, this Capstone Reflection Journal offers a chance to pause, reflect on your learning, and take stock of the challenges and opportunities as you apply your new skills. Below you will find four prompts corresponding with the stages of the *Machine Learning Project Lifecycle*, plus a final, closing reflection.

Upon completion of the capstone, upload this journal in the Capstone Submission Portal, as per the instructions in the course.

SCOPING

Reflection Question: Now that you've reviewed the project plan given to you, how would you explain the 'big picture' goal in your own words? Does the project feel realistic to you? Is there a specific part that looks extra challenging?

DATA AND MODELING

Reflection Question: As you experimented with the model, what settings or features would you consider changing to try to get better results? Did you discover anything messy, missing, or otherwise problematic about the data? If so, how did you fix it so your model could use the information? If you had to explain how this model makes its 'decisions' in simple terms, how would you describe it?

DELIVERY

Reflection Question: How do you know your model actually works? Was it as accurate as you hoped, and if not, what is one thing you would change to make it better?

DEPLOYMENT

Reflection Question: Now that your project is ready for others to use, who do you think will benefit from it the most? What is one thing that might go wrong when a real person starts using the results of the model?

FINAL REFLECTION QUESTIONS FOR COURSE 3

1. You just built a model on real, messy data. Did it behave the way you expected? If not, what did you think was going to happen?
2. The class imbalance was left in on purpose. Did you notice it before you read this, or did your results tip you off? What was your first instinct when you saw the classification report?
3. We picked the Decision Tree because people can follow its logic. But looking at your results, do you think that was the right call for this data, or did we leave performance on the table?

4. Synthetic data, which is artificially generated to mimic real data and is often used to expand dataset size or balance classes, was an option we chose not to pursue. Knowing how the model performs, if you were designing this project from scratch, would you consider using synthetic data? What changes would you make to improve the model or the dataset?
5. Imagine handing this model to someone who has to decide where to spend the next hundred million dollars on carbon capture. What would you tell them it can do, and more importantly, what would you warn them it cannot?
6. What is one thing you would fix if you worked with this data again?



amai.ca | hello@amai.ca

Alberta Machine Intelligence Institute
1101, 10065 Jasper Ave NW
Edmonton, AB T5J 3B1 Canada